

Computer System Organization

Recitation

[Fall 2018]

CSCI-UA 201-006

R09: Multi-thread programming

Recap

- Generating Fibonacci numbers is not parallelizable (unless you use some formulate that does not depend on the previous items).
- A general framework to use threads.
 - Using `pthread_create()` to create threads.
 - The argument that passed to `pthread_create()` may be different for different threads to represent different data.
 - Using `pthread_join()` to wait for all threads to finish.

Multi-thread programming

- If threads have to share data, what should we do?
 - Use lock to protect the data.
 - Use lock-free technique to design the program.
 - ▶ This is an advanced technique, we will only introduce the basic.
- `race.c` demonstrate race condition.


Lock

- In computer science, a lock or mutex (from mutual exclusion) is a synchronization mechanism for enforcing limits on access to a resource in an environment where there are many threads of execution. A lock is designed to enforce a mutual exclusion concurrency control policy.
— Wikipedia
- In short, we use locks to make sure only one thread access the data per time.

Lock

- Coarse-grained lock and fine-grained lock:
 - Coarse-grained lock means that using only one (big) lock to protect the entire data.
 - Fine-grained lock means that using several (small) locks to protect the data. Each lock protect a portion of data.
- Coarse-grained lock is easy to use but has slow performance. On the other hand, fine-grained lock is difficult to use but has much higher performance.

Lock



Lock-free programming

- Problem with locking:
 - Deadlock
 - Fairness
 - Overall performance
 - A lot of problems...
- Lock-free programming
 - Thread-safe access to shared data without the use of synchronization primitives such as mutexes.
 - Possible but not practical in the absence of hardware support.

How I debug multi-threading programs

- Use logging...
- Only use GDB when there is a segmentation fault.

r06 Questions